



JavaOneSM

Sun's 2000 Worldwide Java Developer Conference[®]



Tomcat at Jakarta.Apache.ORG

**Anil Vijendran, Mandar Raje,
Craig McClanahan
Java™ Software,
Sun Microsystems, Inc.**

Session Overview

- Introduction
- Tomcat:
The Servlet Container
- Jasper:
The JSP™ Technology-based Engine
- Future Directions:
The Catalina Architecture



Introduction

- Goals of the Jakarta project
- Projects under the Jakarta umbrella
- Where are we now at Jakarta?



Jakarta Project Goals

- Ensure widespread availability of Java[™] technology-based web technologies, mainly Servlets and JavaServer Pages[™] technology; Tomcat as a reference implementation
- Build a development community around Tomcat and evolve Tomcat into a high quality servlet engine
- Source availability under Apache License to promote easy embedding in commercial products



Projects Under the Jakarta Umbrella

- Tomcat—Servlet container
 - Core support for Servlet 2.2+
 - Jasper 1.1+ engine, for JSP™ technology
 - Native connector(s) to Apache, IIS, NES
- Watchdog—Compliance test suite for the JSP and servlet APIs
- Jakarta Taglibs—a public repository of JSP 1.1 technology-based tags



Projects Under the Jakarta Umbrella

- Ant—Java™ technology-based build tool
- Regexp—Regular expression processing
- Slide—WebDAV support
- Struts—Model-View-Controller (MVC) based framework for web applications



Where Are We Now at Jakarta?

- Tomcat
 - Stable release (04/00): Tomcat 3.1
 - Planned release (07/00): Tomcat 3.2
 - Next major release of Tomcat based on the Catalina architecture
- Key features
 - Support for Java Servlet 2.2 API, JSP 1.1
 - Integrated to run with
 - Apache
 - Microsoft IIS
 - Netscape[™] Enterprise Server





JavaOneSM
Sun's 2000 Worldwide Java Developer Conference

Tomcat: The Servlet Container

Anil Vijendran
(akv@Eng.Sun.COM)
Staff Engineer
Sun Microsystems, Inc.

Design Overview of the Tomcat Servlet Engine

- Brief Intro to Servlets
- Tomcat Design Goals
- Request Processing Overview
- Core Abstractions in Tomcat
- Customizing Request Processing—Interceptors
- Tomcat Configuration



Java™ Servlet API— A Quick Introduction

- Java™ Servlet API (“Java Servlets”)—
Java™ technology-based components
that extend web servers
- Managed by a servlet container
- Have their own lifecycle—init(),
destroy(), service()
- Loaded/unloaded on demand
- Mapped to URL namespace
- Multi-threaded



Tomcat Design Goals

- Main Design Goals
 - Highly customizable—encourages embedding in other app servers, web servers
 - Keep the “core” small, entirely Java technology and portable
 - Stability and compliance over performance and features
- Non-Goals
 - Replacing other web servers
 - Turn Tomcat into a full-fledged application server



Overview of Request Processing in Tomcat

- HTTP server (Apache) receives a request; finds out it is for a servlet or JSP[™] page
- A Connector creates Request/Response objects and calls the ContextManager
 - Connector can use TCP to send the request to the Java technology process or any other method
 - The most common connector is AJP12, used in also in Jserv - with implementations for Apache but also IIS, NES



Overview of Request Processing in Tomcat (Cont.)

- ContextManager starts processing by calling RequestInterceptors to process the request
- Interceptors have the same role as modules in Apache - example: lookup the context, session, do authorization and authentication and find the handler (servlet) for the request
- ContextManager calls the servlet



Core Abstractions in Tomcat

- Connector
- ContextManager
- Container
- Request Interceptors
- Context Interceptors
- Request, Response



Core Abstractions in Tomcat (Cont.)

- Connector
 - Receives requests forwarded from webserver
 - One connector implementation per transport/protocol (example: AJP)
 - Create protocol-specific Request/Response and call ContextManager
- ContextManager
 - Abstract servlet execution/lifecycle from Connector



Core Abstractions in Tomcat (Cont.)

- RequestInterceptor
 - Customize various aspects of request processing—parsing, authorization, authentication, pre/post actions
- ContextInterceptor
 - Notification of various events related to contexts—add/remove web apps, changes to mappings etc



Core Abstractions in Tomcat (Cont.)

- Request/Response
 - Just containers for HTTP request/response state
 - State variables initialized by connectors and interceptors
- Container
 - A group of URLs that share common properties (handler, security constraints)



Customizing Request Processing—Interceptors

```
class UserHomes extends
BaseInterceptor {
    public int contextMap(req, res) {
        uri=req.getRequestURI();
        if (uri.startsWith( "~" ) ) {
            // extract string following ~
            // create context with a
            // docbase=/home/user/webapps
            req.setContext(ctx);
        }
    }
}
```



Tomcat Configuration

- Mostly configured by APIs (to allow embeddability)
- For standalone Tomcat
 - Rule-based XML processor
 - Rules associate components of the XML tree with actions
 - These actions call Tomcat's core APIs to configure Tomcat





Jasper: The JSP™ Technology-based Engine

Mandar Raje

(mandar@Eng.Sun.COM)

Member of Technical Staff

Sun Microsystems, Inc.

A Brief Introduction to JSP™ Technology

- Generation of dynamic content
- Easier way to write servlets
- Separation of logic from presentation
- Use of tags to minimize scripting
- Better tool support
- Web layer of the Java™ 2 Platform, Enterprise Edition (J2EE™)



JSP Technology Example Page

- A simple JSP Technology printing Request parameters:

```
<html>
  <h1> Request Information </h1>
  request URI:<%= request.getURI()%>
  request method:<%= request.getMethod()%>
</html>
```



An Equivalent Servlet

```
public void service (req, res) {  
  
    out.print ("<html>");  
    out.print ("request URI:");  
    out.print (request.getURI());  
    out.print ("request Method:");  
    out.print (request.getMethod());  
    .  
    out.print ("</html>");  
}
```



The Jasper JSP Engine

- Initialization
- Parser and Code Generator
 - Parsing template text
 - Parsing JSP elements
- Compiling generated Servlet
- Debugging
 - Translation-time exceptions
 - Runtime exceptions



Initialization

- Generating a unique class name
 - Avoiding name collisions
 - Using path information in package names
 - Class reloading
- Assigning implicit variables
 - Setting the content type
 - Setting the buffer size
 - Creating context, session (optional)



Parser and Code Generator

- Handling template text
 - Embedded in `out.print()`
- Handling JSP expressions
 - Embedded in `out.print()`
- Handling scriptlets
 - Copied verbatim in the `service()` method
- Handling declarations
 - Define class variables and methods



Parser and Code Generator

- Handling jsp:useBean

```
<jsp:useBean id="foo" scope="request"  
    class="myClass" />
```

```
myClass foo = pageContext.getAttribute  
    ("foo", REQUEST_SCOPE);  
if (foo == null) {  
    foo = (myClass) Beans.Instantiate  
        (loader, "myClass");  
    pageContext.setAttribute  
        ("foo", foo, REQUEST_SCOPE);  
}
```



Parser and Code Generator

- Handling setProperty and getProperty
 - Invoking methods on the beans
- Handling runtime includes and forwards
 - Use underlying Servlet RequestDispatcher mechanism
- Handling custom-tags
 - Parse Tag Library Descriptors (TLDs)
 - Generate code accordingly



Debugging

- Translation-time exceptions
 - Parse-time exceptions are errors in pages

```
<jsp:getProperty nme="foo" .... />
```

error message correctly points the cause and the location where the error has occurred
 - Compile-time exceptions are errors in the generated Java source code (Servlet)

```
<% out.println ("Welcome") %>
```

error message is generated by the Java compiler (javac or jikes)



Debugging

- Runtime exceptions
 - NullPointerException, ArrayOutOfBoundsException
 - <% String foo = null;
if (foo.equals(...)) ... %>
 - Error-page directive works only with RuntimeExceptions
 - <% page errorPage="/error.jsp" %>
 - Error-page element in deployment descriptor



The Jasper Engine


- The Jasper engine is a servlet by design
- Registered using deployment descriptor

```
<servlet>
  <servlet-name> jsp </servlet-name>
  <servlet-class> JspServlet </servlet-class>
</servlet>

<servlet-mapping>
  <servlet-name> jsp </servlet-name>
  <url-pattern> *.jsp </url-pattern>
</servlet-mapping>
```



Jasper Architecture



Compiling Generated Servlet

- **JavaCompiler Interface**
 - Set encoding, classpath, outputdir
 - JikesJavaCompiler, SunJavaCompiler
 - Can be set through init-parameter
- **Handling JSP components with large static data**
 - Limitation on the method size
 - Read static data from a separate file
 - Can be set through init-parameter



Future Directions

- Moving to Catalina
- Better debugging support
- Better tool support
- Generating portable servlets from JSP technology pages
- Support for JSP.next





JavaOneSM
Sun's 2000 Worldwide Java Developer Conference

Future Directions: The Catalina Architecture

Craig McClanahan
(craigmcc@Eng.Sun.COM)
Staff Engineer
Sun Microsystems, Inc.

Diverse Deployment Environments

- Stand alone “mini-application server”
- Connected to an existing web server
- Integrated in full function application server
- Integrated in a development tool
- Embedded in a hardware device
- Multiple communications protocols



Plug-In Functionality Support

- Internal architecture based on components, described by Java interfaces
- Implementation classes configurable at run time (with appropriate defaults)
- Component configuration based on JavaBeans[™] property APIs
- Component lifecycle support and events



Extensible Request Processing

- At various levels (server, host, context, servlet)
- Examples:
 - Security (authentication/access control)
 - Customized logging
 - Resource management
 - Transaction support
- Configurable at runtime



Catalina Component Families

- Connector Family
 - *Request and Response*
- Container Family
 - *Engine, Host, Context, and Wrapper*
- Session Management Family
 - *Manager and Session*
- Utility Components Family
 - *Loader, Logger, Realm, and Resources*



Connector Components

- *Connector*—Represents a communications connection from a remote client
- *Request*—Represents the incoming HTTP request
- *Response*—Represents the outgoing HTTP response



The *Connector* Interface

```
public interface Connector {  
    ...  
    public Container getContainer();  
    public void setContainer(Container container);  
    public Request createRequest();  
    public Response createResponse();  
    ...  
}
```



Container Components

- *Engine*—the entire servlet container
- *Host*—a virtual host
- *Context*—a servlet context
 - Also known as a web application
- *Wrapper*—an individual servlet

All of these components implement the *Container* interface




The *Container* Interface

```
public interface Container {  
    ...  
    public void invoke(Request request,  
        Response response) throws IOException,  
        ServletException;  
    public Container map(Request request,  
        boolean update);  
    ...  
}
```



The Basic Request/Response Processing Model



Adding Custom Request Processing Functionality

- *Valves* can be inserted in the request/response processing paths
- What can a *Valve* do with the request?
 - Examine the request and pass it on
 - Modify the request and pass it on
 - Complete the response and return
- What can a *Valve* do with the response?
 - Examine the response and pass it on
 - Modify the response and pass it on



The *Valve* Interface

```
public interface Container {  
    ...  
    public Valve getNext();  
    public void invoke(Request request,  
        Response response) throws IOException,  
        ServletException;  
    ...  
}
```




Containers Can Implement Pipelines of *Valves*

- *invoke()* method receives a request
- Passes it to the first attached *Valve*
- The first *Valve* can:
 - Preprocess the request and pass it on
 - Complete the request and return
- Each subsequent *Valve* behaves in the same manner
- After processing is complete, the call stack unwinds
 - Local variables to save state
 - You can postprocess the response



Pipelined Request Processing



Session Management Components

- *Manager*—manages all of the sessions for a *Context*
- *Session*—represents an individual **HttpSession**

Session managers can implement persistence, swapping, load balancing, and other features



Utility Components (per Container)

- *Loader*—class loader to use
- *Logger*—implement the `ServletContext.log()` methods
- *Realm*—integrate with external source of users and their roles
- *Resources*—integrate with local filesystem, web server, or application server to retrieve static resources

Containers may inherit utility components from their parent

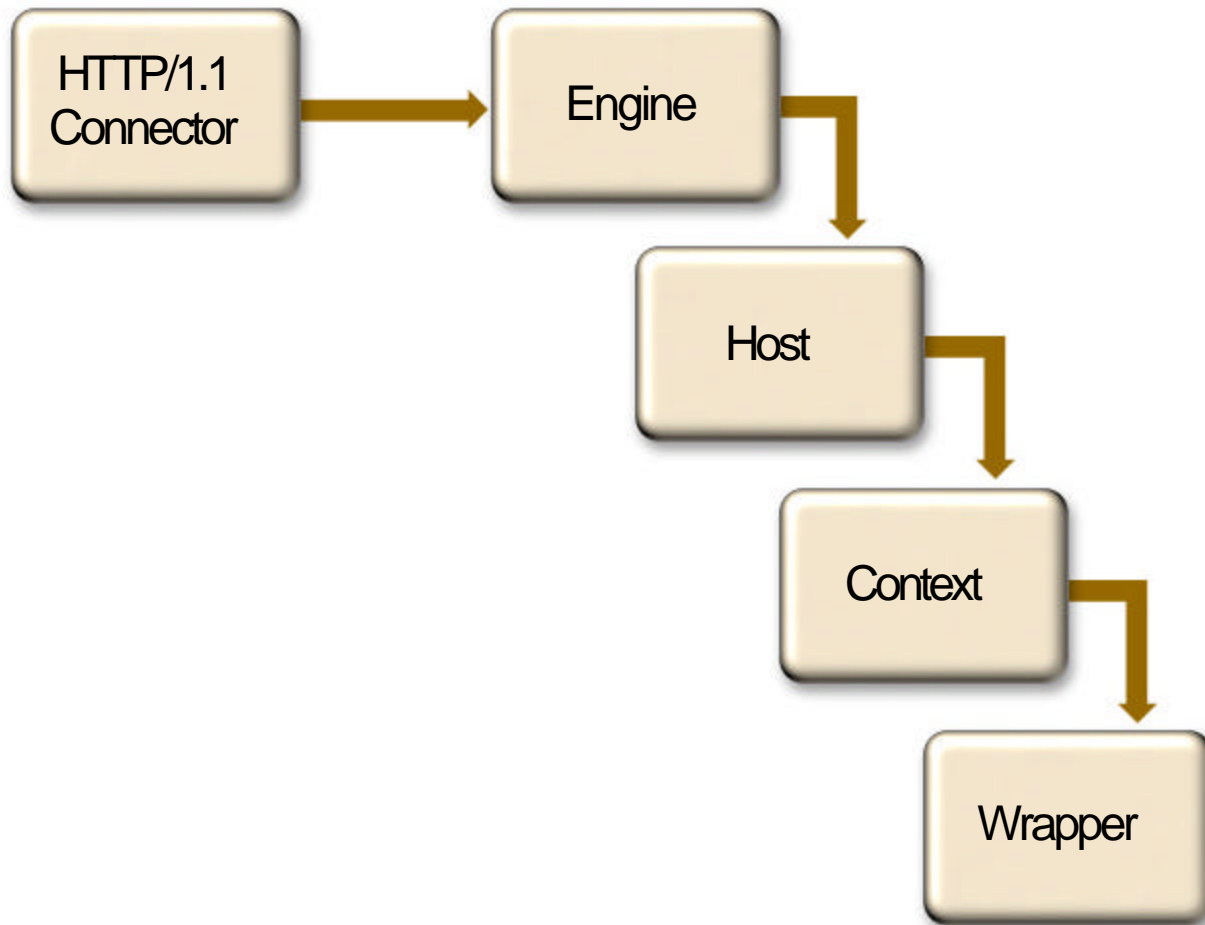


Deployment Scenarios


- Stand-alone server
- Connected to a web server
- Integrated with an application server
- Integrated with a development tool



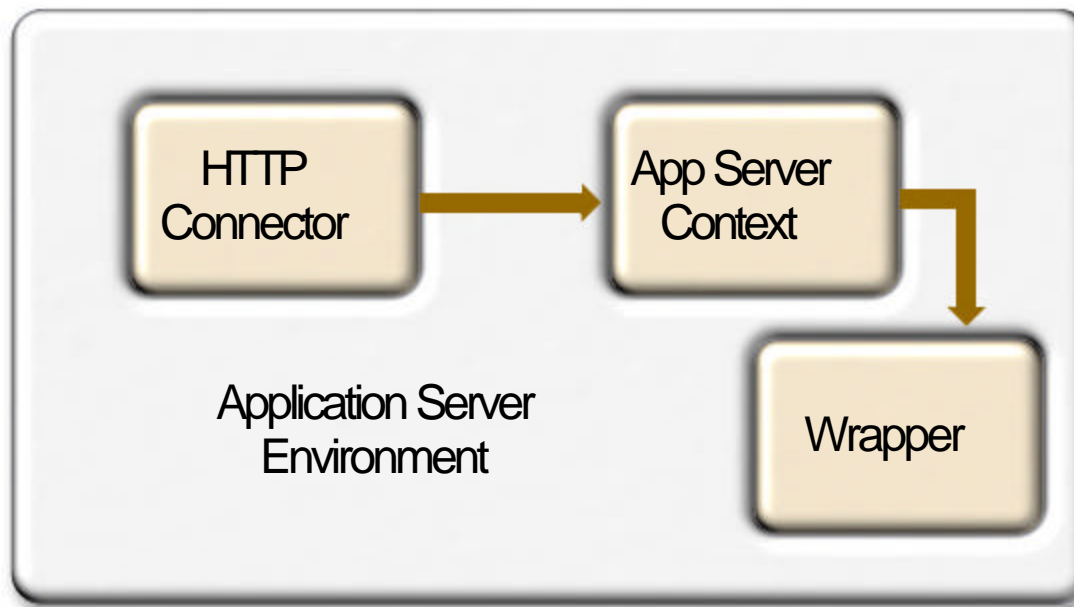
Stand-alone Server



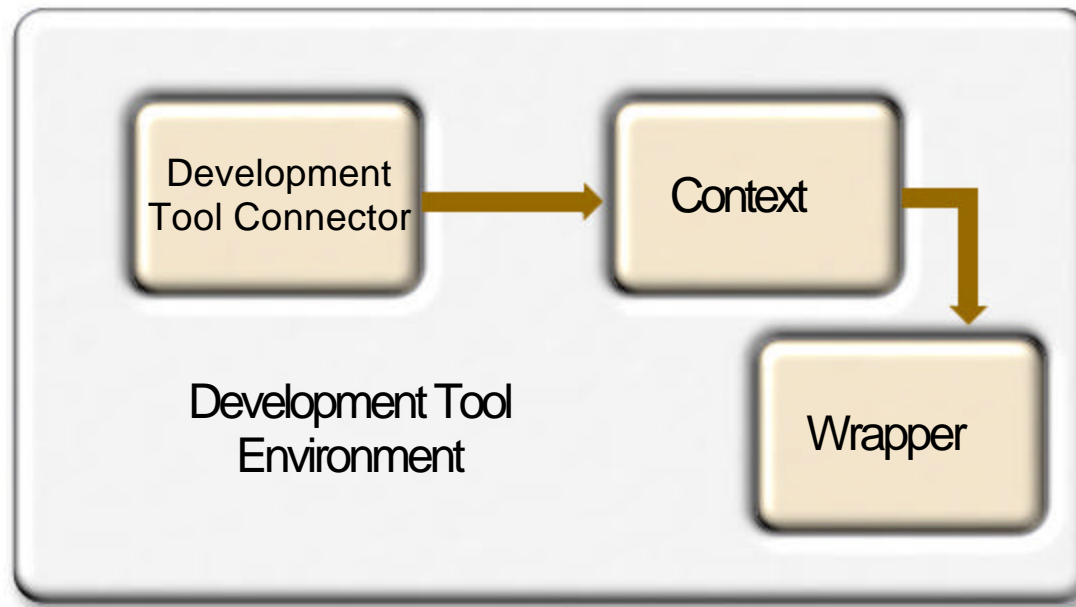
Connected To A Web Server



Integrated with an Application Server



Integrated with a Development Tool



Current Development Status

- Core containers are complete
- HTTP/1.1 connector is complete
- Major remaining functionality needed:
 - Web connectors (Apache, IIS, NES)
 - Distributed server support
 - Administration application
 - SSL/TLS support (standalone mode)
 - Support Servlet 2.3 / JSP 1.2 features
- Testing and tuning is needed



Resources

- Servlet and JSP technology information:
 - <http://java.sun.com/products/servlet>
 - <http://java.sun.com/products/jsp>
 - Mailing lists SERVLET-INTEREST and JSP-INTEREST
- Jakarta Project information
 - <http://jakarta.apache.org>
 - Mailing lists TOMCAT-DEV and TOMCAT-USER





JavaOneSM

Sun's 2000 Worldwide Java Developer Conference*